

Data Science: is a branch of computer science where we study how to store, use and analyze data for deriving information from it.

▼ Pandas

- Pandas is a Python library used for working with data sets.
- It has functions for analyzing, cleaning, exploring, and manipulating data.
- The name "Pandas" has a reference to both "Panel Data", and "Python Data Analysis" and was created by Wes McKinney in 2008.
- Pandas allows us to analyze big data and make conclusions based on statistical theories.
- Pandas can clean messy data sets, and make them readable and relevant.
- Pandas are also able to delete rows that are not relevant, or contains wrong values, like empty or NULL values. This is called cleaning the data.
- Pandas is fast and it has high performance & productivity for users.

Why Use Pandas?

- Fast and efficient for manipulating and analyzing data.
- Data from different file objects can be easily loaded.
- Flexible reshaping and pivoting of data sets
- Provides time-series functionality.

Uses of Pandas:

- Data set cleaning, merging, and joining.
- Easy handling of missing data (represented as NaN) in floating point as well as non-floating point data.
- Columns can be inserted and deleted from DataFrame and higher dimensional objects.
- Powerful group by functionality for performing split-apply-combine operations on data sets.
- Data Visualization

▼ Getting Started

Installing Pandas

- We need to install pandas library using the following **pip command**:

```
pip install pandas
```

Importing Pandas

- After installing pandas on the system, we have to import it before any use, using the following statement:

```
import pandas as pd
```

▼ Pandas Data Structures

Pandas provide following 2 data structures for manipulating data:

1. Series
2. DataFrame

1. Series:

- Pandas Series is a one-dimensional labeled array capable of holding data of any type (integer, string, float, python objects, etc.).
- Labels need not be unique but must be a hashable type.
- In the real world, a Pandas Series will be created by loading the datasets from existing storage, storage can be SQL Database, CSV file, or an Excel file.
- Pandas Series can be created from lists, dictionaries, and from scalar values, etc.

```
import pandas as pd
import numpy as np

arr = np.array([2,3,4,5,6,7,8])
sr = pd.Series(arr)
print(sr)
```

```

0    2
1    3
2    4
3    5
4    6
5    7
6    8
dtype: int64

```

2. DataFrame:

- Pandas DataFrame is a two-dimensional size-mutable, potentially heterogeneous tabular data structure with labeled axes (rows and columns).
- A Data frame is a two-dimensional data structure, i.e., data is aligned in a tabular fashion in rows and columns.
- Applications of DataFrame:
 - Work on Dataset
 - Analysis
 - Dropping
 - Processing
 - Cleaning
 - Join multiple data (CSV, excel format data)
 - Create excel, json, CSV, binary files.
 - Mathematical and Statistical Operations.
 - Use of Group by Function.

Columns: Also called as - Features, variables, field, dimensions.

Rows: Also called as - Records, values, observations, index.

Creating DataFrames:

```

# data = {'a':[1,2,3], 'b':[11,12,13], 'c':[21,22]}
# ValueError: All arrays must be of the same length
# df = pd.DataFrame(data)
# print(df)

```

1. Using a Dictionary with values as lists:

```
data = {'a':[1,2,3], 'b':[11,12,13], 'c':[21,22,23]}
```

```
df = pd.DataFrame(data)
print(df)
```

	a	b	c
0	1	11	21
1	2	12	22
2	3	13	23

2. Can fill same value to all rows:

```
d = {'name': 'Snehal', 'age': 22, 'subjects': ['C', 'C++', 'HTML', 'Java', 'Python']}
df = pd.DataFrame(d)
print(df)
```

	name	age	subjects
0	Snehal	22	C
1	Snehal	22	C++
2	Snehal	22	HTML
3	Snehal	22	Java
4	Snehal	22	Python

3. From dicstionary of numpy arrays:

```
a = np.array([1,2,3,4])
b = np.array(['A', 'B', 'C', 'D'])
c = np.array(['Kop', 'San', 'Sat', 'Pune'])
```

```
d = {'id':a, 'name':b, 'address':c}
df = pd.DataFrame(d)
print(df)
```

	id	name	address
0	1	A	Kop
1	2	B	San
2	3	C	Sat
3	4	D	Pune

4. Create DataFrame from list of lists:

```
lst = [['id', 'name', 'address'], [1,2,3,4], ['A', 'B', 'C', 'D'], ['Kop', 'San', 'Sat', 'Pune']]

df = pd.DataFrame(dict(zip(lst[0], lst[1:])))
print(df)
```

	id	name	address
0	1	A	Kop
1	2	B	San
2	3	C	Sat
3	4	D	Pune

```
print(type(df))  
  
<class 'pandas.core.frame.DataFrame'>
```

▼ Importing and Exporting DataFrame

1. CSV File:

```
# Write to CSV file:  
  
df.to_csv('df_csv.csv')  
  
# Read from CSV file  
ddf = pd.read_csv('df_csv.csv')  
print(ddf)
```

	Unnamed: 0	id	name	address
0	0	1	A	Kop
1	1	2	B	San
2	2	3	C	Sat
3	3	4	D	Pune

2. Excel File:

```
# Write to Excel File  
df.to_excel('df_xl.xlsx')  
  
# Read from Excel File  
dex = pd.read_excel('df_xl.xlsx')  
print(dex)
```

	Unnamed: 0	id	name	address
0	0	1	A	Kop
1	1	2	B	San
2	2	3	C	Sat
3	3	4	D	Pune

3. JSON File:

```
# Write to json file:  
df.to_json('df_json.json')
```

```
# Read from json file:
dj = pd.read_json('df_json.json')
print(dj)
```

```
   id name address
0    1    A      Kop
1    2    B      San
2    3    C      Sat
3    4    D      Pune
```

4. HTML File:

```
# Write to HTML File
df.to_html('df_html.html')
```

```
# Read from html file:
dh = pd.read_html('df_html.html')
print(dh)
```

```
[   Unnamed: 0  id name address
0           0    1    A      Kop
1           1    2    B      San
2           2    3    C      Sat
3           3    4    D      Pune]
```

▼ DataFrame Functions

Check size of data frame:

```
df = pd.read_csv('Housing.csv')
print(df)
```

```
   price  area  bedrooms  bathrooms  stories  mainroad  guestroom  basement  \
0  13300000  7420         4          2         3        yes         no         no
1  12250000  8960         4          4         4        yes         no         no
2  12250000  9960         3          2         2        yes         no         yes
3  12215000  7500         4          2         2        yes         no         yes
4  11410000  7420         4          1         2        yes         yes        yes
..    ...    ...         ...        ...        ...        ...        ...        ...
540  1820000  3000         2          1         1        yes         no         yes
541  1767150  2400         3          1         1        no         no         no
542  1750000  3620         2          1         1        yes         no         no
543  1750000  2910         3          1         1        no         no         no
544  1750000  3850         3          1         2        yes         no         no

   hotwaterheating  airconditioning  parking  prefarea  furnishingstatus
0                no                yes         2        yes        furnished
1                no                yes         3        no        furnished
```

```

2          no          no          2          yes  semi-furnished
3          no          yes          3          yes    furnished
4          no          yes          2          no    furnished
..         ...         ...         ...         ...         ...
540         no          no          2          no    unfurnished
541         no          no          0          no  semi-furnished
542         no          no          0          no    unfurnished
543         no          no          0          no    furnished
544         no          no          0          no    unfurnished

```

```
[545 rows x 13 columns]
```

```
7085
```

```
RangeIndex(start=0, stop=545, step=1)
```

```
print(df.size) # --> rows * cols
```

```
7085
```

```
print(df.index) # --> Range of index from Start to End
```

```
RangeIndex(start=0, stop=545, step=1)
```

Get Names of the Columns:

```
print(df.columns) # --> Names of columns
```

```

Index(['price', 'area', 'bedrooms', 'bathrooms', 'stories', 'mainroad',
       'guestroom', 'basement', 'hotwaterheating', 'airconditioning',
       'parking', 'prefarea', 'furnishingstatus'],
      dtype='object')

```

```
print(df.axes) # --> Range and Names of the columns
```

```

[RangeIndex(start=0, stop=545, step=1), Index(['price', 'area', 'bedrooms', 'bathroom',
       'guestroom', 'basement', 'hotwaterheating', 'airconditioning',
       'parking', 'prefarea', 'furnishingstatus'],
      dtype='object')]

```



df.info():

- Getting info of overall data frame.

```
print(df.info())
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 13 columns):
#   Column          Non-Null Count  Dtype
---  -

```

```

0    price      545 non-null    int64
1    area       545 non-null    int64
2    bedrooms   545 non-null    int64
3    bathrooms  545 non-null    int64
4    stories    545 non-null    int64
5    mainroad   545 non-null    object
6    guestroom  545 non-null    object
7    basement   545 non-null    object
8    hotwaterheating 545 non-null object
9    airconditioning 545 non-null object
10   parking    545 non-null    int64
11   prefarea   545 non-null    object
12   furnishingstatus 545 non-null object
dtypes: int64(6), object(7)
memory usage: 55.5+ KB
None

```

df.describe():

- Return all the statistical functions values.
- For all the columns with numeric data type, present in the dataframe.
- Does not work for the string(object) data type.

```
print(df.describe())
```

```

           price      area  bedrooms  bathrooms  stories  \
count  5.450000e+02    545.000000  545.000000  545.000000  545.000000
mean    4.766729e+06    5150.541284    2.965138    1.286239    1.805505
std     1.870440e+06    2170.141023    0.738064    0.502470    0.867492
min     1.750000e+06    1650.000000    1.000000    1.000000    1.000000
25%     3.430000e+06    3600.000000    2.000000    1.000000    1.000000
50%     4.340000e+06    4600.000000    3.000000    1.000000    2.000000
75%     5.740000e+06    6360.000000    3.000000    2.000000    2.000000
max     1.330000e+07   16200.000000    6.000000    4.000000    4.000000

           parking
count  545.000000
mean    0.693578
std     0.861586
min     0.000000
25%     0.000000
50%     0.000000
75%     1.000000
max     3.000000

```

df.head():

- Return first 5 rows, by default.
- Can also specify number of rows to fetch.

```
print(df.head())
```


	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	\
0	13300000	7420	4	2	3	yes	no	no	
1	12250000	8960	4	4	4	yes	no	no	
2	12250000	9960	3	2	2	yes	no	yes	
3	12215000	7500	4	2	2	yes	no	yes	
4	11410000	7420	4	1	2	yes	yes	yes	

	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
0	no	yes	2	yes	furnished
1	no	yes	3	no	furnished
2	no	no	2	yes	semi-furnished
3	no	yes	3	yes	furnished
4	no	yes	2	no	furnished

```
print(df.head(10))
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	\
0	13300000	7420	4	2	3	yes	no	no	
1	12250000	8960	4	4	4	yes	no	no	
2	12250000	9960	3	2	2	yes	no	yes	
3	12215000	7500	4	2	2	yes	no	yes	
4	11410000	7420	4	1	2	yes	yes	yes	
5	10850000	7500	3	3	1	yes	no	yes	
6	10150000	8580	4	3	4	yes	no	no	
7	10150000	16200	5	3	2	yes	no	no	
8	9870000	8100	4	1	2	yes	yes	yes	
9	9800000	5750	3	2	4	yes	yes	no	

	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
0	no	yes	2	yes	furnished
1	no	yes	3	no	furnished
2	no	no	2	yes	semi-furnished
3	no	yes	3	yes	furnished
4	no	yes	2	no	furnished
5	no	yes	2	yes	semi-furnished
6	no	yes	2	yes	semi-furnished
7	no	no	0	no	unfurnished
8	no	yes	2	yes	furnished
9	no	yes	1	yes	unfurnished

df.tail():

- Return last 5 rows, by default.
- Can also specify number of rows to fetch.

```
print(df.tail())
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	\
540	1820000	3000	2	1	1	yes	no	yes	
541	1767150	2400	3	1	1	no	no	no	
542	1750000	3620	2	1	1	yes	no	no	
543	1750000	2910	3	1	1	no	no	no	
544	1750000	3850	3	1	2	yes	no	no	

	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
540	no	no	2	no	unfurnished
541	no	no	0	no	semi-furnished
542	no	no	0	no	unfurnished
543	no	no	0	no	furnished
544	no	no	0	no	unfurnished

```
print(df.tail(10))
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	\
535	2100000	3360	2	1	1	yes	no	no	
536	1960000	3420	5	1	2	no	no	no	
537	1890000	1700	3	1	2	yes	no	no	
538	1890000	3649	2	1	1	yes	no	no	
539	1855000	2990	2	1	1	no	no	no	
540	1820000	3000	2	1	1	yes	no	yes	
541	1767150	2400	3	1	1	no	no	no	
542	1750000	3620	2	1	1	yes	no	no	
543	1750000	2910	3	1	1	no	no	no	
544	1750000	3850	3	1	2	yes	no	no	

	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
535	no	no	1	no	unfurnished
536	no	no	0	no	unfurnished
537	no	no	0	no	unfurnished
538	no	no	0	no	unfurnished
539	no	no	1	no	unfurnished
540	no	no	2	no	unfurnished
541	no	no	0	no	semi-furnished
542	no	no	0	no	unfurnished
543	no	no	0	no	furnished
544	no	no	0	no	unfurnished

isna():

- Show null values.
- Return DataFrame which contains -
 - True - for NULL values
 - False - for NON-NULL values.

```
d = {'id':[1,2,3], 'name':['A','B','C'], 'age':[21,23,np.nan]}
```

```
df = pd.DataFrame(d)
df.isna()
```

	id	name	age
0	False	False	False
1	False	False	False
2	False	False	True

Transpose of the DataFrame:

- Convert the row indices to column names and vice versa.

```
df.iloc[1:5,:5].transpose()
```

	1	2
id	2	3
name	B	C
age	23.0	NaN

```
df.iloc[1:5,:5].T
```

	1	2
id	2	3
name	B	C
age	23.0	NaN

▼ Accessing DataFrame

Access by name of the Column:

- Can access using the following 2 methods:

```
1. df.col_name
```

```
2. df['col_name']
```

```
df = pd.read_csv('Housing.csv')
```

```
print(df.price)
```

```
0    13300000
1    12250000
2    12250000
3    12215000
4    11410000
...
```

```

540      1820000
541      1767150
542      1750000
543      1750000
544      1750000
Name: price, Length: 545, dtype: int64

```

```

# Give column name as index in []
print(df.price[0])

```

```
13300000
```

```
print(df['price'])
```

```

0      13300000
1      12250000
2      12250000
3      12215000
4      11410000
...
540     1820000
541     1767150
542     1750000
543     1750000
544     1750000
Name: price, Length: 545, dtype: int64

```

```

# Give name of the column as index in []
print(df['price'][34])

```

```
8120000
```

df.loc:

- It also access the actual values at the index and columns.
- Can get record at an index:

```
df.loc[index]
```

```
df.loc[start:end]
```

```
df.loc[start:end:step]
```

```
df.loc[10]
```

```

price      9800000
area       13200
bedrooms    3
bathrooms   1

```

```

stories                2
mainroad               yes
guestroom              no
basement               yes
hotwaterheating        no
airconditioning         yes
parking                2
prefarea               yes
furnishingstatus        furnished
Name: 10, dtype: object

```

```
df.loc[10:12]
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotw
10	9800000	13200	3	1	2	yes	no	yes	
11	9681000	6000	4	3	2	yes	yes	yes	
12	9310000	6550	4	2	2	yes	no	no	



```
df.loc[10:22:2]
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotw
10	9800000	13200	3	1	2	yes	no	yes	
12	9310000	6550	4	2	2	yes	no	no	
14	9240000	7800	3	2	2	yes	no	no	
16	9100000	6600	4	2	2	yes	yes	yes	
18	8890000	4600	3	2	2	yes	yes	no	
20	8750000	4320	3	1	2	yes	no	yes	
22	8645000	8050	3	1	1	yes	yes	yes	



- Can access records with a condition in **loc[]**:

```
df.loc[condition]
```

```
df.loc[df['bedrooms']==3]
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	ho
2	12250000	9960	3	2	2	yes	no	yes	
5	10850000	7500	3	3	1	yes	no	yes	
9	9800000	5750	3	2	4	yes	yes	no	
10	9800000	13200	3	1	2	yes	no	yes	
14	9240000	7800	3	2	2	yes	no	no	
...
533	2100000	2400	3	1	2	yes	no	no	
537	1890000	1700	3	1	2	yes	no	no	
541	1767150	2400	3	1	1	no	no	no	
543	1750000	2910	3	1	1	no	no	no	
544	1750000	3850	3	1	2	yes	no	no	


300 rows × 13 columns



```
df.loc[df['bedrooms']<3]
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hot
61	7070000	8880	2	1	1	yes	no	no	
66	6930000	13200	2	1	1	yes	no	yes	
73	6685000	6600	2	2	4	yes	no	yes	
91	6419000	6750	2	1	1	yes	yes	yes	
114	6020000	6800	2	1	1	yes	yes	yes	
...
535	2100000	3360	2	1	1	yes	no	no	
538	1890000	3649	2	1	1	yes	no	no	
539	1855000	2990	2	1	1	no	no	no	
540	1820000	3000	2	1	1	yes	no	yes	
542	1750000	3620	2	1	1	yes	no	no	

138 rows × 13 columns



Access Specific Value in the Data Frame:

```
df[col_name][index]
```

```
df['price'][23]
```

```
8645000
```

```
df['price'][23:45]
```

```
23    8645000
24    8575000
25    8540000
26    8463000
27    8400000
28    8400000
29    8400000
30    8400000
31    8400000
32    8295000
33    8190000
34    8120000
35    8080940
36    8043000
37    7980000
38    7962500
39    7910000
40    7875000
41    7840000
42    7700000
43    7700000
44    7560000
```

```
Name: price, dtype: int64
```

Access Data from Multiple Columns:

- Use names of the columns in the form of a list:

```
df[[col_names_list]]
```

```
df[['price', 'area', 'bathrooms']]
```

iloc:

- Pass axes numbers for index and columns.

```
df.iloc[1:5,]  
# Return all columns of 1 to 4 rows
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwa
1	12250000	8960	4	4	4	yes	no	no	
2	12250000	9960	3	2	2	yes	no	yes	
3	12215000	7500	4	2	2	yes	no	yes	
4	11410000	7420	4	1	2	yes	yes	yes	



```
df.iloc[1:5,:3]  
# Return only firsy 3 columns of rows 1:4.
```

	price	area	bedrooms	
1	12250000	8960	4	
2	12250000	9960	3	
3	12215000	7500	4	
4	11410000	7420	4	

[Colab paid products](#) - [Cancel contracts here](#)