

```
import numpy as np
import pandas as pd
import datetime
```

▼ Feature Engineering

- Feature Engineering is the process of transforming data to increase the predictive performance of machine learning models.

Data normalization:

- Data Normalization could also be a typical practice in machine learning which consists of transforming numeric columns to a standard scale.


```
data = {
    'col1': [1,2,3,4,5],
    'col2': [8,7,3,6,4],
    'col3': [12000, 23000, 45000, 34000, 21000]
}
df = pd.DataFrame(data)
```

```
print(df)
```

```
   col1  col2  col3
0     1     8  12000
1     2     7  23000
2     3     3  45000
3     4     6  34000
4     5     4  21000
```

```
import matplotlib.pyplot as plt
df.plot(kind = 'bar')
# When plotting the graph, we can see that,
# col3 has dominated the graph and we cannot observe the others.
```

☐ <Axes: >




```
# So, we can scale all column data in smaller ranges
for column in df.columns:
    df[column] = df[column] / df[column].abs().max()

print(df)
# Now, we can see all the col values are closer to each other
```

```
   col1  col2  col3
0   0.2  1.000  0.266667
1   0.4  0.875  0.511111
2   0.6  0.375  1.000000
3   0.8  0.750  0.755556
4   1.0  0.500  0.466667
```

```
import matplotlib.pyplot as plt
df.plot(kind = 'bar')
# Now, we can observe features from the graph.
```



Scaling:

- In cases where all the columns have a significant difference in their scales, are needed to be modified in such a way that all those values fall into the same scale. This process is called **Scaling**.
- There are two most common techniques of how to scale columns of Pandas dataframe –
 1. Min-Max Normalization
 2. Standardization.

```
df = pd.read_csv('IRIS.csv')
```

```
print(df.head())
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

1. min-max Normalization:

- Here, all the values are scaled in between the range of [0,1] where 0 is the minimum value and 1 is the maximum value.
- The formula for Min-Max Normalization is –

$$x_{\text{norm}} = (x - x_{\text{min}}) / (x_{\text{max}} - x_{\text{min}})$$

```
print(df.info())
```

```
# We can perform scaling on numeric data cols only, so for scaling purpose,
# have to drop the last col with "species" for this operation only
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   sepal_length    150 non-null   float64
1   sepal_width     150 non-null   float64
2   petal_length    150 non-null   float64
3   petal_width     150 non-null   float64
4   species         150 non-null   object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
None
```

```
new_df = df.drop('species', axis=1) # Have to use axis=1, for column
```

```
# Now, perform scaling on each value of each column
df_norm = ((new_df-new_df.min()) / (new_df.max()-new_df.min()))

# Now, merge the normalized df with the dropped species col
final_df = pd.concat((df_norm, df.species), axis=1)

print(final_df)
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	0.222222	0.625000	0.067797	0.041667	Iris-setosa
1	0.166667	0.416667	0.067797	0.041667	Iris-setosa
2	0.111111	0.500000	0.050847	0.041667	Iris-setosa
3	0.083333	0.458333	0.084746	0.041667	Iris-setosa
4	0.194444	0.666667	0.067797	0.041667	Iris-setosa
..
145	0.666667	0.416667	0.711864	0.916667	Iris-virginica
146	0.555556	0.208333	0.677966	0.750000	Iris-virginica
147	0.611111	0.416667	0.711864	0.791667	Iris-virginica
148	0.527778	0.583333	0.745763	0.916667	Iris-virginica
149	0.444444	0.416667	0.694915	0.708333	Iris-virginica

[150 rows x 5 columns]

```
# Using MinMaxScaler from sklearn
```

```
from sklearn.preprocessing import MinMaxScaler
import pandas as pd

# Drop species col, having str data
new_df = df.drop('species', axis=1)

scaler = MinMaxScaler()

df_scaled = scaler.fit_transform(new_df.to_numpy())
df_scaled = pd.DataFrame(df_scaled, columns=[
    'sepal_length', 'sepal_width', 'petal_length', 'petal_width'])

print("Scaled Dataset Using MinMaxScaler")
df_scaled.head()
```

Scaled Dataset Using MinMaxScaler

	sepal_length	sepal_width	petal_length	petal_width
0	0.222222	0.625000	0.067797	0.041667
1	0.166667	0.416667	0.067797	0.041667
2	0.111111	0.500000	0.050847	0.041667
3	0.083333	0.458333	0.084746	0.041667
4	0.194444	0.666667	0.067797	0.041667

2. Standardization:

- Standardization doesn't have any fixed minimum or maximum value.
- Here, the values of all the columns are scaled in such a way that they all have a mean equal to 0 and standard deviation equal to 1.
- This scaling technique works well with outliers.
- Thus, this technique is preferred if outliers are present in the dataset.

```
import pandas as pd
from sklearn.preprocessing import StandardScaler

# Drop species col, having str data
new_df = df.drop('species', axis=1)

std_scaler = StandardScaler()

df_scaled = std_scaler.fit_transform(new_df.to_numpy())
df_scaled = pd.DataFrame(df_scaled, columns=[
    'sepal_length', 'sepal_width', 'petal_length', 'petal_width'])

print("Scaled Dataset Using StandardScaler")
df_scaled.head()
```

Scaled Dataset Using StandardScaler

	sepal_length	sepal_width	petal_length	petal_width
0	-0.900681	1.032057	-1.341272	-1.312977
1	-1.143017	-0.124958	-1.341272	-1.312977
2	-1.385353	0.337848	-1.398138	-1.312977



Time Series Analysis and Resampling

- Pandas provide a different set of tools using which we can perform all the necessary tasks on date-time data.

Working with datetime data:

#1: Create a dates dataframe

```
# Create dates dataframe with frequency
data = pd.date_range('1/1/2011', periods = 15, freq = 'M')
```

data

```
DatetimeIndex(['2011-01-31', '2011-02-28', '2011-03-31', '2011-04-30',
                '2011-05-31', '2011-06-30', '2011-07-31', '2011-08-31',
                '2011-09-30', '2011-10-31', '2011-11-30', '2011-12-31',
                '2012-01-31', '2012-02-29', '2012-03-31'],
              dtype='datetime64[ns]', freq='M')
```

2: Current date and time

```
# x = pd.datetime.now()
# x.date, x.month, x.year
```

```
stamp = pd.Timestamp(datetime.datetime(2023, 8, 4))
stamp = pd.Timestamp(datetime.datetime.now())
print(stamp)
```

```
result = stamp.today()
print(result)
print(stamp.day_name())
print(stamp.day_of_week)
print(stamp.weekday())
print(stamp.hour)
print(stamp.day)
print(stamp.days_in_month)
print(stamp.date())
print(stamp.dayofyear)
print(stamp.daysinmonth)
print(stamp.minute)
```

```
2023-08-04 14:25:05.287975
2023-08-04 14:25:05.289076
Friday
4
4
14
4
31
2023-08-04
216
31
25
```

Create date and time with dataframe



```
rng = pd.DataFrame()
rng['date'] = pd.date_range(datetime.datetime.now(), periods = 72, freq = 'H')
```

```
# Print the dates in dd-mm-yy format
rng[:5]
```

Create features for year, month, day, hour, and minute

```
rng['year'] = rng['date'].dt.year
rng['month'] = rng['date'].dt.month
rng['day'] = rng['date'].dt.day
rng['hour'] = rng['date'].dt.hour
rng['minute'] = rng['date'].dt.minute
rng['second'] = rng['date'].dt.second
```

```
# Print the dates divided into features
rng.head(3)
```

		date	year	month	day	hour	minute	second		
0	2023-08-04 14:25:09.473671	2023		8	4	14	25	9		
1	2023-08-04 15:25:09.473671	2023		8	4	15	25	9		
2	2023-08-04 16:25:09.473671	2023		8	4	16	25	9		

Resampling and time-based indexing:

Handling time zones and date offsets:

- Dateoffsets are a standard kind of date increment used for a date range in Pandas.
- DateOffsets can be created to move dates forward a given number of valid dates.
- Pandas `tseries.offsets.DateOffset` is used to create standard kind of date increment used for a date range.
- **Syntax:**

```
pandas.tseries.offsets.DateOffset(n=1, normalize=False, **kwds)
```

Parameters:

- `n` : The number of time periods the offset represents.
- `normalize` : Whether to round the result of a DateOffset addition down to the previous midnight.
- `level` : int, str, default None
- `**kwds` : Temporal parameter that add to or replace the offset value. Parameters that add to the offset (like `Timedelta`): years, months etc.
- Returns : DateOffsets

```
# Creating Timestamp
ts = pd.Timestamp('2019-10-10 07:15:11')
print(ts)

# Create the DateOffset
do = pd.tseries.offsets.DateOffset(n = 2)
print(do)

# We can now add the DateOffset to any date, to increment the date
new_date = ts + do
print(new_date)

# Thus, the date would move forward by 2 days.
```

```
2019-10-10 07:15:11
<2 * DateOffsets>
2019-10-12 07:15:11
```

```
# Providing additional arguments to the DateOffset

today = pd.Timestamp(datetime.datetime.now())

# Create new DateOffset
date_off = pd.tseries.offsets.DateOffset(days=10, hours=3, minutes=15)
# Move the timestamp by 10 days, 3 hrs and 15 mins

future_date = today + date_off

print(today)
print(future_date)

2023-08-04 14:25:21.270023
2023-08-14 17:40:21.270023
```

Working with datetime data in pandas:

Datetime features can be divided into two categories:

- The first one: time moments in a period

- Second: the time passed since a particular period.
- These features can be very useful to understand the patterns in the data.

Divide a given date into features –

- `pandas.Series.dt.year` returns the year of the date time.
- `pandas.Series.dt.month` returns the month of the date time.
- `pandas.Series.dt.day` returns the day of the date time.
- `pandas.Series.dt.hour` returns the hour of the date time.
- `pandas.Series.dt.minute` returns the minute of the date time.

```
#3: Break date and time into separate features

# Create date and time with dataframe
rng = pd.DataFrame()
rng['date'] = pd.date_range(datetime.datetime.now().date(), periods = 72, freq = 'H')

# Print the dates in dd-mm-yy format
rng[:5]

# Create features for year, month, day, hour, and minute
rng['year'] = rng['date'].dt.year
rng['month'] = rng['date'].dt.month
rng['day'] = rng['date'].dt.day
rng['hour'] = rng['date'].dt.hour
rng['minute'] = rng['date'].dt.minute

# Print the dates divided into features
rng.head(10)
```

	date	year	month	day	hour	minute
0	2023-08-04 00:00:00	2023	8	4	0	0
1	2023-08-04 01:00:00	2023	8	4	1	0
2	2023-08-04 02:00:00	2023	8	4	2	0
3	2023-08-04 03:00:00	2023	8	4	3	0
4	2023-08-04 04:00:00	2023	8	4	4	0
5	2023-08-04 05:00:00	2023	8	4	5	0
6	2023-08-04 06:00:00	2023	8	4	6	0
7	2023-08-04 07:00:00	2023	8	4	7	0
8	2023-08-04 08:00:00	2023	8	4	8	0
9	2023-08-04 09:00:00	2023	8	4	9	0

```
# Import dataset
data = pd.read_csv('COVID_19_Containment_measures_data.csv')

# print(data.loc[0])

sample = {'id': data['ID'], 'country': data['Country'], 'start': data['Date Start'],
          'end': data['Date end intended']}
df = pd.DataFrame(sample)

# print(df)
print(data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1703 entries, 0 to 1702
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0    ID                                    820 non-null    object
1    Applies To                           29 non-null     object
2    Country                              1675 non-null   object
3    Date Start                           1639 non-null   object
4    Date end intended                    242 non-null    object
5    Description of measure implemented  1640 non-null   object
6    Exceptions                           41 non-null     object
7    Implementing City                    127 non-null    object
8    Implementing State/Province          179 non-null    object
9    Keywords                            1615 non-null   object
10   Quantity                             302 non-null    float64
11   Source                              1517 non-null   object
12   Target city                          1 non-null      object
```

```

13 Target country      132 non-null    object
14 Target region      29 non-null    object
15 Target state        0 non-null    float64
dtypes: float64(2), object(14)
memory usage: 213.0+ KB
None

```

```

# Convert the Time column to datetime format
df['start'] = pd.to_datetime(df.start)
df['end'] = pd.to_datetime(df.end)

```

```
print(df.head())
```

```

   id  country  start end
0  163  Austria 2020-03-16 NaT
1  132  Germany 2020-02-01 NaT
2  578  United Kingdom 2020-03-20 NaT
3  372  United Kingdom 2020-03-16 NaT
4  357  United Kingdom 2020-03-16 NaT

```

```
df.dtypes
```

```
# Thus, shows that, start and end columns are converted to datetime type
```

```

id          object
country     object
start       datetime64[ns]
end         datetime64[ns]
dtype: object

```

```
# Get details from the date type from the DF
```

```

# Get hour detail from time data
print(df.start.dt.hour.head())

```

```

# Get name of each date
# df.start.dt.weekday_name.head()

```

```

# Get ordinal day of the year
print(df.start.dt.dayofyear.head())

```

```

0    0.0
1    0.0
2    0.0
3    0.0
4    0.0
Name: start, dtype: float64
0    76.0
1    32.0
2    80.0
3    76.0
4    76.0
Name: start, dtype: float64

```

Resampling time series data (e.g., downsampling and upsampling):

- Pandas **dataframe.resample()** function is primarily used for time series data.
- A time series is a series of data points indexed (or listed or graphed) in time order.
- It is a Convenience method for frequency conversion and resampling of time series.
- Object must have a datetime-like index (DatetimeIndex, PeriodIndex, or TimedeltaIndex), or pass datetime-like values to the on or level keyword.

Syntax :

```
DataFrame.resample(rule, how=None, axis=0, fill_method=None, closed=None, label=None, convention='start', kind=None, lof
```

Parameters :

- **rule** : the offset string or object representing target conversion
- **axis** : int, optional, default 0
- **closed** : {'right', 'left'}
- **label** : {'right', 'left'}
- **convention** : For PeriodIndex only, controls whether to use the start or end of rule

- **loffset** : Adjust the resampled time labels
 - **base** : For frequencies that evenly subdivide 1 day, the "origin" of the aggregated intervals. For example, for '5min' frequency, base could range from 0 through 4. Defaults to 0.
 - **on** : For a DataFrame, column to use instead of index for resampling. Column must be datetime-like.
 - **level** : For a MultiIndex, level (name or number) to use for resampling. Level must be datetime-like.
- Resampling generates a unique sampling distribution on the basis of the actual data.
 - Most commonly used time series frequency are –
 - **W** : weekly frequency
 - **M** : month end frequency
 - **SM** : semi-month end frequency (15th and end of month)
 - **Q** : quarter end frequency

```
data = pd.read_csv('COVID_19_Containment_measures_data.csv')
sample = {'id': data['ID'], 'country': data['Country'], 'start_date': data['Date Start'],
          'end': data['Date end intended'], 'quantity': data['Quantity']}
covid_df = pd.DataFrame(sample)
covid_df.dropna(inplace=True)
covid_df.to_csv('covid_data.csv')
```

```
# By default the "start_date" column was in string format,
# we need to convert it into date-time format
```

```
# parse_dates=["start_date"], converts the "start_date"
# column to date-time format. We know that
# resampling works with time-series data only
# so convert "start_date" column to index
```

```
# index_col="start_date", makes "start_date" column, the index of the data frame
df = pd.read_csv('covid_data.csv', parse_dates=["start_date"], index_col="start_date")
```

```
# Printing the dataframe
df
```

	Unnamed: 0	id	country	end	quantity
start_date					
2020-03-16	378	254	Finland	Apr 13, 2020	10.0
2020-03-16	414	292	Lithuania	Mar 30, 2020	50.0
2020-03-13	452	345	Latvia	Apr 14, 2020	200.0
2020-03-15	453	342	Netherlands	Apr 06, 2020	30.0
2020-03-15	468	364	Ireland	Mar 29, 2020	10.0
2020-03-16	503	396	United States	Mar 30, 2020	1000.0
2020-03-08	529	447	Romania	Mar 31, 2020	1000.0
2020-03-11	530	436	Romania	Mar 31, 2020	100.0
2020-03-12	533	442	Norway	Mar 26, 2020	50.0
2020-03-13	544	459	Germany	Apr 20, 2020	90.0
2020-03-10	599	588	Russia	Apr 10, 2020	5000.0
2020-03-16	613	583	Russia	Apr 10, 2020	50.0
2020-03-13	615	574	Russia	Mar 20, 2020	1000.0
2020-03-16	632	618	US: Oregon	Apr 16, 2020	25.0
2020-03-18	940	861	Denmark	Apr 13, 2020	10.0

```
# Resampling the time series data based on weekly frequency
# we apply it on stock open price 'W' indicates week
weekly_resampled_data = df.quantity.resample('W').mean()
```

```
# find the mean opening price of each week
# for each week over the period
weekly_resampled_data
```

```
start_date
2020-03-08    1000.000000
2020-03-15     810.000000
```



```
2020-03-22    190.833333
Freq: W-SUN, Name: quantity, dtype: float64
```

Working with time-based data in pandas (e.g., datetime index):

```
# index_col ="start_date", makes "start_date" column, the index of the data frame
df = pd.read_csv('covid_data.csv', parse_dates=["start_date"], index_col ="start_date")

df['end'] = pd.to_datetime(df.end)
# Printing the dataframe
df
```

Unnamed: 0	id	country	end	quantity
start_date				
2020-03-16	378 254	Finland	2020-04-13	10.0
2020-03-16	414 292	Lithuania	2020-03-30	50.0
2020-03-13	452 345	Latvia	2020-04-14	200.0
2020-03-15	453 342	Netherlands	2020-04-06	30.0
2020-03-15	468 364	Ireland	2020-03-29	10.0
2020-03-16	503 396	United States	2020-03-30	1000.0
2020-03-08	529 447	Romania	2020-03-31	1000.0
2020-03-11	530 436	Romania	2020-03-31	100.0
2020-03-12	533 442	Norway	2020-03-26	50.0
2020-03-13	544 459	Germany	2020-04-20	90.0
2020-03-10	599 588	Russia	2020-04-10	5000.0
2020-03-16	613 583	Russia	2020-04-10	50.0
2020-03-13	615 574	Russia	2020-03-20	1000.0
2020-03-16	632 618	US: Oregon	2020-04-16	25.0
2020-03-18	940 861	Denmark	2020-04-13	10.0